

B551 - Assignment 5

Authors: Anup Bharadwaj, Raghuveer Kanchibail Krishnamurthy, Supreeth Keragodu
Suryaprakash

Neural Network:

A brief description of how the program works, assumptions and formulation of the problem is provided as a comment on top of 'nnet.py' file.

We tried varying some of the parameters of the neural network to see if they improve in accuracies. Following are some of the results observed with varying parameters:

Size of training set:

The accuracies increased greatly with increase in the size of the training set initially and slowly the difference became smaller. We couldn't record results of all the tests we made, but below are some of the results observed:

No. of images in the training set	Accuracy (%)	Running time (hh:mm:ss)
40000	64.36	00:23:40
12000	63.62	00:04:40
2000	55.99	00:00:47

Learning rate:

Learning rate had the most impact in accuracies for training a neural network. The accuracies dropped if the learning rate was too high or too low. Since I was using a smooth approximation for rectifier functions, the values would overflow out of the supported range for most learning rates. Following are some of the results observed for varied learning rate that successfully trained the model:

Learning rate	Accuracy (%)	Running time (hh:mm:ss)
1/length of training data	55 – 65 (depends on training data)	00:00:47 – 00:23:40
0.001	59.59 (12000 files)	00:06:39
0.1	49.2 (12000 files)	00:13:37
0.00001	62.77 (12000 files)	00:05:21

Activation Function:

I experimented on 4 Activation functions – ReLu, Leaky ReLu, smooth approximation of ReLu and sigmoid functions. The best performance was by leaky reLu and the smooth approximation of ReLu. ReLu performed worst as most outputs would come to zero and the network stopped converging. Below are the results when ran for the full training dataset:

Activation Function	Accuracy (%)	Running time (hh:mm:ss)
Leaky ReLu	64.26	00:23:29
Smooth approximation of ReLu	64.36	00:23:40

ReLu	~25% (converged to a single orientation)	~00:10:00 (Lost the output file)
Sigmoid	61.86%	~ 00:23:00 (Lost the output file)

Traditional vs stochastic gradient:

Both gave very similar results but traditional way took longer to train the model. Since the initial weights are random, there were slight differences in the accuracies observed. The only significant difference noticed was the time taken. Stochastic gradient was able to train 40000 images in under 30 minutes while traditional approach took over an hour.

Initial weight distribution:

Since, we assign the initial weights randomly, we tried two ways of assigning initial weights – pseudorandom generator and Gaussian. The range did matter as higher ranges resulted in overflow errors and Gaussian distribution of randomness yielded better accuracies.

Random generator	Accuracies	Time taken
Pseudo-random generator	~53-61	~00:23:00
Gaussian	~62-64	~00:23:00

Sample run

Some of the classifications of a sample run is as follows:

File name	True Orientation	Prediction
13962200152	180	0
434740832	90	90
2304830497	180	180
5430983827	180	90
4413965131	0	180

In most cases, the model got confused with inverted pictures, the errors were mostly misclassification of 0 as 180 (vice versa) and 90 as 270 (and vice versa). Full results can be found at the file: sample_test in github.

Nearest Neighbors:

<i>Value of K</i>	Accuracy in %	Run Time in Mins(Approx)
<i>1</i>	67.24	19.86
<i>7</i>	68.14	20.86
<i>8</i>	69.78	21.12
<i>9</i>	69.78	21.46
<i>10</i>	72.33	21.78

In Nearest neighbor the performance varies a lot depending upon the size of training set and the size of the test set. The run time grows exponentially with increase in the dataset. This is not at all recommended for huge training/test data. The memory usage will also be very high.

Sample run

Some of the classifications of a sample run is as follows:

File name	True Orientation	Prediction
10351347465	270	180
10684428096	90	270
12178947064	90	270
22272865	270	270
22728648085	90	90

The model predicted most of the images with orientation 270 and 90 in a wrong way. This is the pattern we could observe in this model.

ADABOOST:

<i>Number of Stumps</i>	<i>Accuracy %</i>	<i>Run Time in Secs(Approx)</i>
10	36.37	11.52
25	38.17	14.77
50	39.76	21.93
100	47.29	34.47
200	54.5	53.41
500	62.63	121.40

Run time is directly proportional to the size of training data. The Accuracy increases with the increase in Decision Stumps.

Sample run

Some of the classifications of a sample run is as follows:

File name	True Orientation	Prediction
10008707066	0	0
10107730656	180	180
10814107565	180	0
10577249185	180	0
10161556064	0	270

The model predicted the inverted images in a wrong way. This is the pattern we could observe.

Recommendation:

We would recommend using a Neural net classifier with Leaky ReLu as activation function, $1/\text{len}(\text{dataset})$ as the learning rate, 64 hidden nodes and a Gaussian random value generator to converge faster. Even though the smooth approximation of ReLu provides better accuracies, it is vulnerable to overflows unless handled correctly.

While nearest neighbor performs better, it is very susceptible to overfitting and takes a long time to predict in production.